Fujiao Ji

Ph.D in Computer Science, University of Tennessee, Knoxville
fji1@vols.utk.edu | (865) 824-7093 | fujiaoji.github.io
Expected Graduation Date: May 2027

RESEARCH INTERESTS

Web Security; Security Benchmarking; Agentic AI Security.

EDUCATION & EXPERIENCE

University of Tennessee, Knoxville Aug. 2022 – Present
Ph.D. in Computer Science, Advisor: Dr. Doowon Kim Knoxville, TN, USA

- Graduate Research / Teaching Assistant: Research on phishing website measurement and detection.
- Coursework: Computer Systems Organization, Algorithms, Deep Learning, Machine Learning

Shandong University of Science and Technology Aug. 2018 – Jun. 2021
M.S. in Computer Software and Engineering, Advisor: Prof. Zhongying Zhao Qingdao, China

- Graduate Research / Teaching Assistant. Research on heterogeneous graph representation learning.

Baidu Map Jun. 2021 – Jan. 2022
Research Intern, Advisor: Dr. Yibo Sun and Dr. Lei Shao Beijing, China

- Named entity recognition for Chinese address data.

SELECTED PUBLICATIONS

- PhishGen: Understanding LLM Capabilities in Automated Phishing Website Generation.
  Fujiao Ji, Yerim Kim, Yena Cho, Seonhye Park, Hyoungshick Kim, and Doowon Kim. *arXiv*: 2610.30226, 2026.
- How Can We Effectively Use LLMs for Phishing Detection?: Evaluating the Effectiveness of Large Language Model-based Phishing Detection Models.
  Fujiao Ji and Doowon Kim. *arXiv*: 2511.09606, 2025.
- Evaluating the effectiveness and robustness of visual similarity-based phishing detection models.
  Fujiao Ji, Kiho Lee, Hyungjoon Koo, Wenhao You, Euijin Choo, Hyoungshick Kim, and Doowon Kim. *USENIX Security Symposium*, 2025.
- What's in Phishers: A Longitudinal Study of Security Configurations in Phishing Websites and Kits.
  Kyungchan Lim, Kiho Lee, **Fujiao Ji**, Yonghwi Kwon, Hyoungshick Kim, and Doowon Kim. *Web Conference*, 2025.
- Occasionally secure: A comparative analysis of code generation assistants. Elgedawy, Ran, Porter Dosch, John Sadik, Senjuti Dutta, Anuj Gautam, Konstantinos Georgiou, Farzin Gholamrezae, **Fujiao Ji**, et al. *arXiv*: 2402.00689, 2024.
- A comparative study on heterogeneous information network embeddings.
  Fujiao Ji, Zhongying Zhao, Hui Zhou, Heng Chi, and Chao Li. *Journal of Intelligent & Fuzzy Systems*, 2020.
- Identifying high influential users in social media by analyzing users' behaviors.
  Zhao, Zhongying, Hui Zhou, Bijun Zhang, **Fujiao Ji**, and Chao Li. *Journal of Intelligent & Fuzzy Systems*, 2019.

SELECTED PROJECTS

Agentic AI Security: Working On Jun. 2026 - Present

- Evaluated web-browsing agents during task execution, uncovering gaps between belief and actual browsing behaviors. Identified failures in which agents omitted website security checks and directed users to phishing sites.

Blue-Team Evaluation of LLM-Generated Phishing Websites Jun. 2025 - Aug. 2026

- Designed *PhishGen*, an adversarial evaluation framework spanning 14 commercial, open-source, and uncensored LLMs, 7 multimodal prompt configurations, and 12 prompt-rewriting strategies; generated and analyzed 30,192 phishing website samples across four major brands.
- Built an end-to-end testing pipeline (Python, Playwright, HTML/CSS analysis, automated visual metrics, and browser-based credential submission) to measure safeguard bypass, visual fidelity, frontend rendering, backend execution, and credential persistence; found 96.2% of generated pages rendered successfully while only 38.7% completed the full credential-harvesting workflow.
- Identified novel model- and generator-level security risks, including near-perfect DOM imitation, systematic CSS styling omissions, and a strong textual-over-visual brand mimicry bias; translated these findings into actionable signals for detecting and mitigating LLM-generated phishing attacks.

Security Evaluation of LLM-Based Phishing Detection Jan. 2025 - May 2026

- Built *PhishBench*, a multimodal benchmark of 9,567 real-world phishing websites (APWG eCX, Apr.–Oct. 2025) and 400 manually verified benign sites across 200 brands; ran the controlled comparison of 4 LLMs (GPT, Gemini, Qwen, Llama) against 2 DL baselines over 4 input configurations.
- Found Gemini delivers the best trade-off (89.7% detection at a 2.0% FPR) and GPT the most conservative (73.7% at 0.5%), while open-source LLMs incurred up to 48.0% FPR (Qwen) and DL baselines, though far faster (1.1 s/sample), missed most phishing pages (FNR up to 78.1%).
- Conducted failure analysis revealing that GPT appears anchored to a few salient signals, Gemini responds broadly across Normal-looking features, Qwen shows sparse and potentially unstable cue use, and Llama appears less visually accompanied and more reliant on textual-structural signals.

Adversarial Evaluation of Visual Anti-Phishing Systems Aug. 2022 - Dec. 2024

- Conducted a large-scale adversarial evaluation of visual anti-phishing models using 451K real-world phishing websites collected through APWG.
- Analyzed detection failures and identified three recurring evasion strategies–pipeline exploitation, benign-logo mimicry, and simple visual manipulation–revealing a 20.7% average performance degradation from curated benchmarks to real-world phishing data.
- Designed an adversarial robustness benchmark of 7,223 screenshots spanning 14 visible and five perturbation-based logo manipulations; uncovered transferable vulnerabilities and derived defenses based on adversarial augmentation, OCR–vision integration, multi-cue ensembles, and image preprocessing.

Chinese Address Parsing Project Jun. 2021 - Jan. 2022

- Extracted and structured address data in the format of province, city, district, town, and point of interest.
- Recognized named entities through a biaffine attention mechanism based on the pre-trained model ERNIE 1.0 under the framework of PaddlePaddle and improved performance via post-processing.
- Evaluated the performance of point of interest chunks, where the F1 score is 81.25% for 1,000 real-world data from Baidu Map and 80.41% for 2,985 public data from Chinese Address Corpus.

SKILLS

- **Languages & Frameworks**: Python, PyTorch, TensorFlow, PaddlePaddle, C++
- **Skills**: Deep Learning, RAG, CTF, AI Agents, LLM Distillation, LLM Fine-tuning, Phishing Website Detection, Linux, Tmux, vLLM, Hugging Face, Azure, LaTeX, Pandas, Graph Representation Learning, Vibe Coding, Claude Code, Codex

PROFESSIONAL ACTIVITIES

- Paper Reviewer for USENIX 2027, WWW 2025 and 2027, SANER 2027, Pattern Recognition
- Artifact Evaluation Committee for NDSS 2027, USENIX 2026, IEEE S&P 2026